

# CS 580

# ALGORITHM DESIGN AND ANALYSIS

## Greedy Algorithms 2: Scheduling (cf. KT 4.1 & 4.2)


Vassilis Zikas

# SO FAR

- Shortest paths
  - Dijkstra
  - Bellman-Ford
- Today:
  - Greedy algorithms that don't have to do with graphs!
  - Interval scheduling (4.1 in KT)
  - Scheduling to minimize lateness (4.2 in KT)


# INTERVAL SCHEDULING

- There's an incoming set of jobs  $\{1, \dots, n\}$
- The  $i$ th job corresponds to an interval  $[s_i, t_i]$
- Two jobs are compatible if they don't overlap
- Goal: find maximum subset of compatible jobs




# INTERVAL SCHEDULING

- There's an incoming set of jobs  $\{1, \dots, n\}$
- The  $i$ th job corresponds to an interval  $[s_i, t_i]$
- Two jobs are compatible if they don't overlap
- Goal: find maximum subset of compatible jobs




# INTERVAL SCHEDULING

- There's an incoming set of jobs  $\{1, \dots, n\}$
- The  $i$ th job corresponds to an interval  $[s_i, t_i]$
- Two jobs are compatible if they don't overlap
- Goal: find maximum subset of compatible jobs



# INTERVAL SCHEDULING

- There's an incoming set of jobs  $\{1, \dots, n\}$
- The  $i$ th job corresponds to an interval  $[s_i, t_i]$
- Two jobs are compatible if they don't overlap
- Goal: find maximum subset of compatible jobs



# INTERVAL SCHEDULING GREEDY TEMPLATES

- Earliest start time?
- Earliest finish time?
- Pick smallest size first?
- Minimize number of overlaps?


# INTERVAL SCHEDULING GREEDY TEMPLATES

- Earliest start time?
- Earliest finish time?
- Pick smallest size first?
- Minimize number of overlaps?



# INTERVAL SCHEDULING GREEDY TEMPLATES

- Earliest start time?
- Earliest finish time?
- Pick smallest size first?
- Minimize number of overlaps?



# INTERVAL SCHEDULING: EARLIEST FINISH TIME?

- Consider algorithms in increasing order of finish time  $t_i$ 
    - Take a job if compatible with the ones picked so far
1. Sort by finish time to get  $t_1 \leq t_2 \leq \dots \leq t_n$
  2.  $A \leftarrow \emptyset$ .     //Set of jobs selected so far
  3. For  $j = 1, \dots, n$ :
    - If  $j$  compatible with  $A$ :  $A = A \cup \{j\}$

# INTERVAL SCHEDULING

- Running time  $O(n \log n)$ 
  - Sorting takes  $O(n \log n)$
  - Checking if a job is compatible equivalent to  $s_i \geq f_{j^*}$ , the last job added:  $O(n)$

1. Sort by finish time to get  $t_1 \leq t_2 \leq \dots \leq t_n$
2.  $A \leftarrow \emptyset$ .     //Set of jobs selected so far
3. For  $j = 1, \dots, n$ :
  - If  $j$  compatible with  $A$ :  $A = A \cup \{j\}$

# INTERVAL SCHEDULING

- Theorem: The greedy algorithm is optimal

## Proof

- Assume greedy is not optimal
- Let  $i_1, \dots, i_k$  be the jobs selected by greedy
- Let  $j_1, \dots, j_m$  be the optimal set of jobs
  - $m > k$

By definition,  $i_{r+1}$  finishes before  $j_{r+1}$




# INTERVAL SCHEDULING

- Theorem: The greedy algorithm is optimal

## Proof

- Assume greedy is not optimal
- Let  $i_1, \dots, i_k$  be the jobs selected by greedy
- Let  $j_1, \dots, j_m$  be the optimal set of jobs
  - $m > k$

Replacing  $j_{r+1}$  with  $i_{r+1}$  won't break feasibility



# INTERVAL SCHEDULING

- Theorem: The greedy algorithm is optimal

## Proof

- Assume greedy is not optimal
- Let  $i_1, \dots, i_k$  be the jobs selected by greedy
- Let  $j_1, \dots, j_m$  be the optimal set of jobs
  - $m > k$

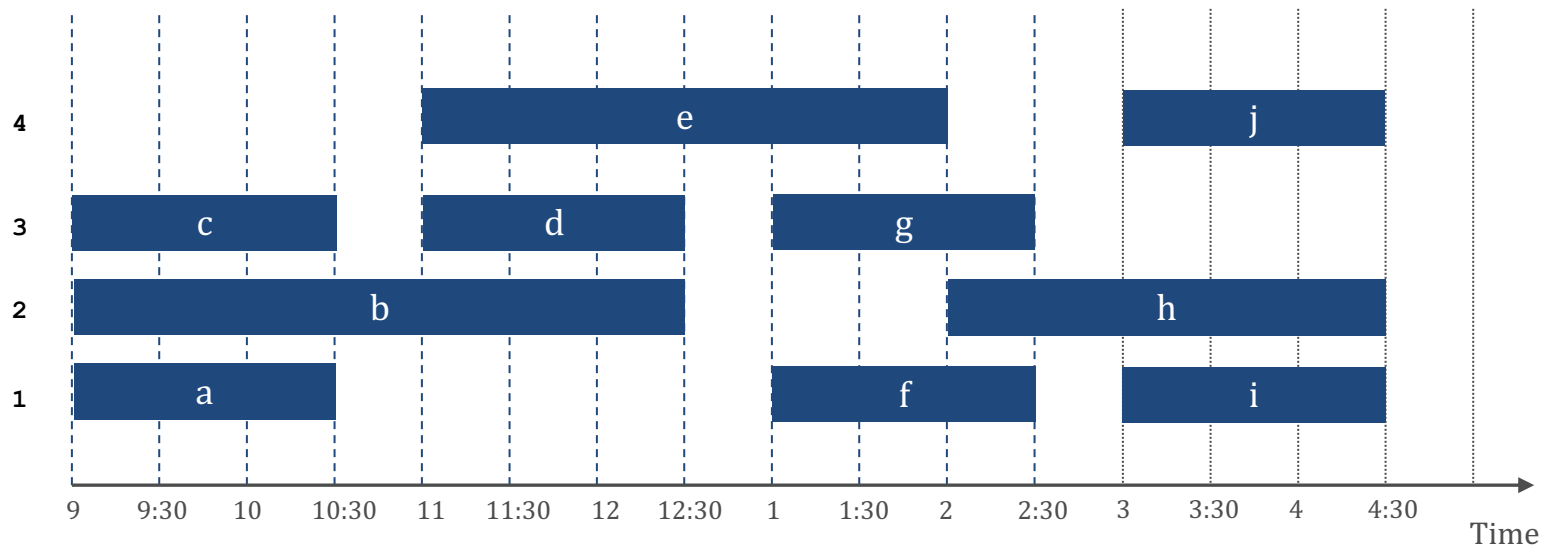
Size of OPT didn't change either  
Contradiction:  $r$  was supposed to be maximal



# INTERVAL PARTITIONING (SCHEDULING ALL INTERVALS)


- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$
- Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room

# INTERVAL PARTITIONING (SCHEDULING ALL INTERVALS)






# INTERVAL PARTITIONING (SCHEDULING ALL INTERVALS)




# INTERVAL PARTITIONING (SCHEDULING ALL INTERVALS)

- Perhaps possible with two?
  - No:  $a$ ,  $b$  and  $c$  overlap, so we need at least 3
- Definition: The **depth** of a set of intervals is the maximum number that pass over any single point in time



# INTERVAL PARTITIONING (SCHEDULING ALL INTERVALS)

- Perhaps possible with two?
  - No:  $a$ ,  $b$  and  $c$  overlap, so we need at least 3
- Definition: The **depth** of a set of intervals is the maximum number that pass over any single point in time



# INTERVAL PARTITIONING

- Question: does there exist a schedule equal to the depth?
  - If the answer is yes, then this is clearly optimal

# INTERVAL PARTITIONING

- Greedy algorithm: consider lectures in increasing starting time
    - assign to any compatible classroom (including a new one)
1. Sort intervals by start time to get  $s_1 \leq \dots \leq s_n$
  2. For  $j = 1$  to  $n$ :
    - If lecture  $j$  is compatible with some classroom  $k$ , schedule  $j$  in  $k$
    - Otherwise, schedule  $j$  in a new classroom

# INTERVAL PARTITIONING

- Implementation:
  - $O(n \log(n))$  to sort
  - At each time, need to figure out if a lecture is compatible
    - Use priority queue to remember end times of classrooms!

# INTERVAL PARTITIONING

**Theorem:** Greedy is optimal

Proof

- Let  $d$  be the number of classrooms greedy needs
- Classroom  $d$  opened because we needed to schedule some job  $j$  that was incompatible with all  $d - 1$  previous classrooms
- These  $d$  jobs (the  $d - 1$  plus  $j$ ) all end after  $s_j$
- Since we sorted by start time, all these jobs also **start** before (or at)  $s_j$
- Therefore, there exists a set of  $d$  jobs that overlap at  $s_j \rightarrow$  need at least  $d$  classrooms


# SCHEDULING TO MINIMIZE LATENESS

- Problem:
  - We have a single resource and need to process  $n$  jobs
  - We can process only one job at a time
  - Job  $j$  requires  $t_j$  units of time and expires at time  $d_j$ 
    - If  $j$  starts at  $s_j$  it finishes at  $f_j = s_j + t_j$
  - Lateness:  $\ell_j = \max\{0, f_j - d_j\}$
  - Goal: minimize maximum latency  $L = \max_j \ell_j$



# SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
$t_j$	3	2	1	4	3	2
$d_j$	6	8	9	9	14	15

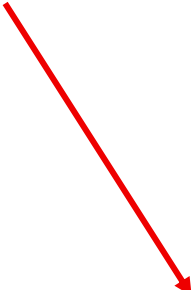


# SCHEDULING TO MINIMIZE LATENESS

- Greedy template: Consider jobs in *some* order
  - Shortest processing time first
    - Sort by  $t_j$  (ascending)
  - Smallest slack first
    - Sort by  $d_j - t_j$  (ascending)

# SCHEDULING TO MINIMIZE LATENESS

- Greedy template: Consider jobs in *some* order
  - Shortest processing time first
    - Sort by  $t_j$  (ascending)
  - Smallest slack first
    - Sort by  $d_j - t_j$  (ascending)

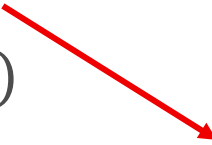


	1	2
$t_j$	1	10
$d_j$	100	10

Processing time  $\neq$  Urgency  
Deadline can be very far in the future

# SCHEDULING TO MINIMIZE LATENESS

- Greedy template: Consider jobs in *some* order
  - Shortest processing time first
    - Sort by  $t_j$  (ascending)
  - Smallest slack first
    - Sort by  $d_j - t_j$  (ascending)



	1	2
$t_j$	1	10
$d_j$	2	10

Deadline - Processing time  $\neq$  Urgency  
Ignores small deadlines that take little time

# SCHEDULING TO MINIMIZE LATENESS


- Greedy algorithm: Earliest deadline first
  - Initially seems like a bad idea to not take processing time into account...

1. Sort jobs so that  $d_1 \leq d_2 \leq \dots \leq d_n$
2.  $t=0$
3. For  $j = 1 \dots n$ :
  - Assign job  $j$  to interval  $[t, t + t_j]$
  - $t = t + t_j$

# SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
$t_j$	3	2	1	4	3	2
$d_j$	6	8	9	9	14	15

max lateness = 1




# SCHEDULING TO MINIMIZE LATENESS

- **Observation:** There exists an optimal schedule with no idle time
  - Proof:
    - Assume the optimal schedule has some time interval  $[t, t']$  where nothing is scheduled
    - For all jobs that start after  $t$ , move their start time earlier by  $t' - t$
    - Still feasible, and latencies only decreased
- **Observation:** Greedy has no idle time

# SCHEDULING TO MINIMIZE LATENESS


- **Definition:** Given a schedule  $S$ , an **inversion** is a pair of jobs  $i$  and  $j$  such that  $d_i < d_j$  but  $j$  is scheduled before  $i$ 
  - By definition, greedy has no inversions





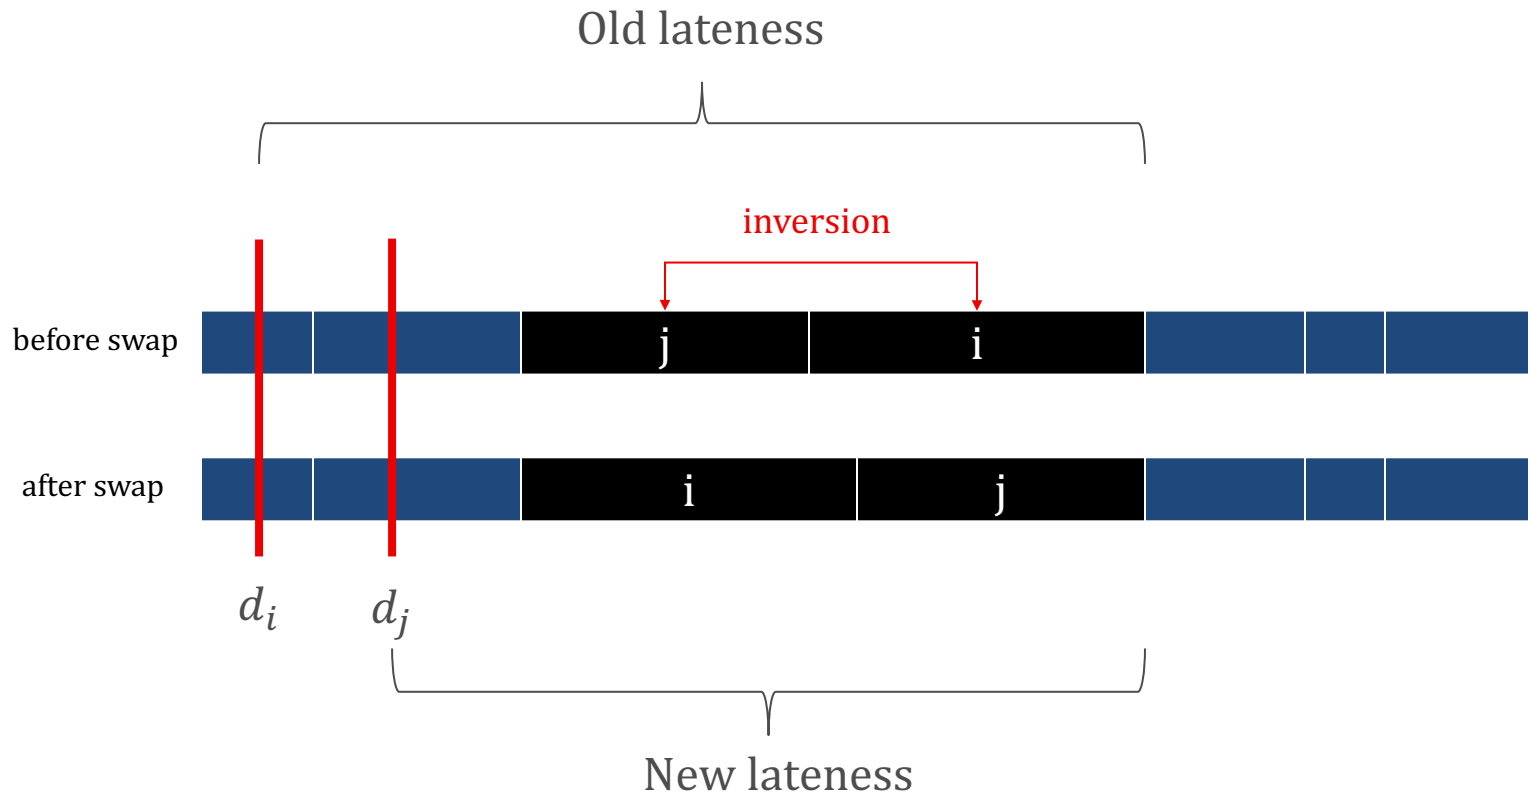
# SCHEDULING TO MINIMIZE LATENESS

- Observation: If an inversion exists, then a “consecutive” inversion exists
  - Start from the left, until you find the first violation




# SCHEDULING TO MINIMIZE LATENESS

- What happens if we swap?
  - Everyone other than  $i$  and  $j$  remain unchanged




# SCHEDULING TO MINIMIZE LATENESS

- What happens if we swap?
  - $\ell'_j$




# SCHEDULING TO MINIMIZE LATENESS

- What happens if we swap?
  - $\ell'_j = f'_j - d_j$




# SCHEDULING TO MINIMIZE LATENESS

- What happens if we swap?
  - $\ell'_j = f'_j - d_j = f_i - d_j$




# SCHEDULING TO MINIMIZE LATENESS

- What happens if we swap?
  - $\ell'_j = f'_j - d_j = f_i - d_j < f_i - d_i$



# SCHEDULING TO MINIMIZE LATENESS

- What happens if we swap?
  - $\ell'_j = f'_j - d_j = f_i - d_j < f_i - d_i = \ell_i$



# SCHEDULING TO MINIMIZE LATENESS

- Theorem: Greedy is optimal
- Proof:
- Let  $S^*$  be an optimal schedule (strictly) better than greedy, with the fewest number of inversions
- Wlog  $S^*$  has no idle times
- If  $S^*$  has no inversions, then  $S^* = S$
- If  $S^*$  has an inversion, let  $(i, j)$  be an adjacent inversion: swap it!
  - Maximum latency didn't decrease (otherwise  $S^*$  is not optimal)
  - But, number of inversions decreased! Contradiction



# SUMMARY

- Interval scheduling
- Interval Partitioning
- Scheduling to minimize latency
- Meta lesson: similar (same?) analysis
  - Gradually transform the solution of any other algorithm to one of greedy, without changing the performance