

CS 580

ALGORITHM DESIGN AND ANALYSIS

Approximation Algorithms: Part 1

Vassilis Zikas


SO FAR

- A wild problem appears
- So far:
 1. Find a polynomial time algorithm 😊
 2. Prove NP-completeness and give up 😞
- But, what if I really need to solve the problem (even if it is NP-hard)?

PICK YOUR POISON


- Three features:
 - Generality — small variations easy to move to
 - Optimality — correctness!
 - Polynomial time
- Pick two!

PICK YOUR POISON

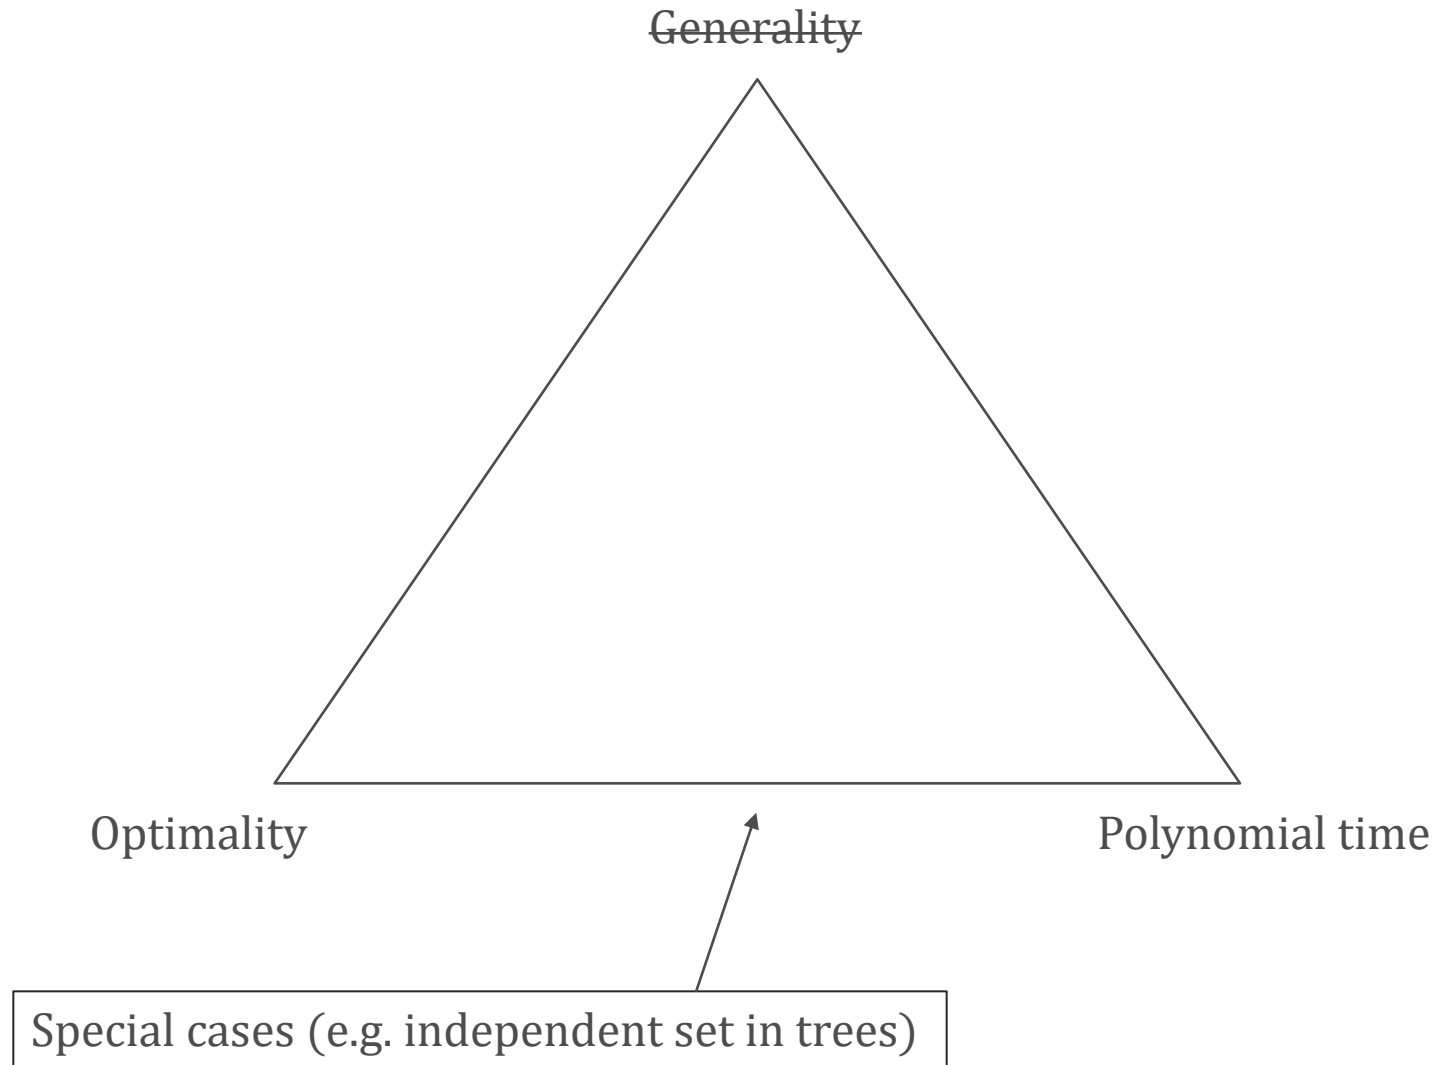


PICK YOUR POISON

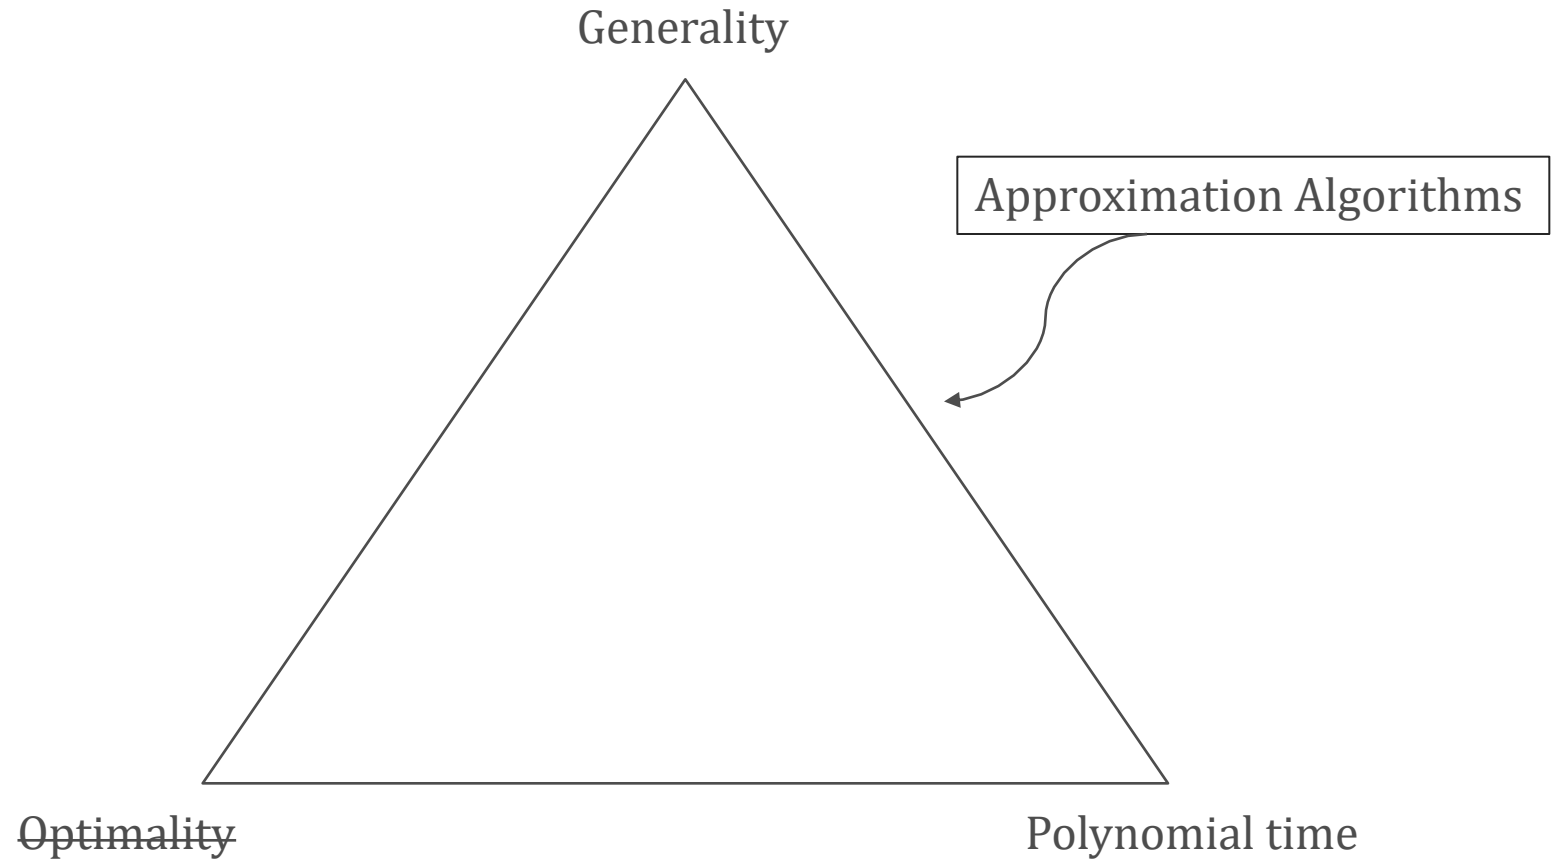
eg gen all permutations!



PICK YOUR POISON



PICK YOUR POISON



APPROXIMATION ALGORITHMS

- ρ -approximation algorithm
 - Runs in polynomial time
 - The ratio of the true optimum to the output of the algorithm is within ρ

APPROXIMATION ALGORITHMS

- E.g., a 2-approximation algorithm for INDEPENDENT SET would return an independent set whose size is **at least** $\frac{1}{2}$ the size of the maximum independent set (maximization problems)
- A 2-approximation algorithm for VERTEX COVER would return a vertex cover whose size is **at most** twice the size of the minimum vertex cover (minimization problems)

TODAY

- LOAD BALANCE (11.1)
- SET COVER (11.3)
- VERTEX COVER

LOAD BALANCING

- Input:
 - m machines: M_1, \dots, M_m
 - n jobs: job j has processing time t_j
- Output:
 - Assign jobs to machines so that machine loads are “balanced”
 - Let $A(i)$ be the set of jobs assigned to M_i
 - Load of M_i is $T_i = \sum_{j \in A(i)} t_j$
 - Goal: Minimize the **makespan** $T = \max_i T_i$

LOAD BALANCING

normally NP complete

- Theorem: LOAD-BALANCING is NP-complete
 - We will not show this, but notice that $m = 2$ is essentially SUBSET SUM
- Our interest:
 - Find, in polynomial time, an assignment with good makespan



LOAD BALANCING

- Greedy algorithm:
 - Take an arbitrary job j
 - Assign it to the machine with the smallest load
 - Repeat

```
List-Scheduling( $m, n, t_1, t_2, \dots, t_n$ ) {  
  for  $i = 1$  to  $m$  {  
     $L_i \leftarrow 0$             $\leftarrow$  load on machine  $i$   
     $A(i) \leftarrow \phi$         $\leftarrow$  jobs assigned to machine  $i$   
  }  
  
  for  $j = 1$  to  $n$  {  
     $i = \operatorname{argmin}_k L_k$   
     $A(i) \leftarrow A(i) \cup \{j\}$   
     $L_i \leftarrow L_i + t_j$   
  }  
  return  $A(1), \dots, A(m)$   
}
```

LOAD BALANCING

- Theorem [Graham 1966!]: **Greedy** is a 2-approximation
- Lemma 1: The optimal makespan T^* is at least $\max_i t_i$
- Proof:
 - Some machine must process the largest job

LOAD BALANCING

- Lemma 2: The optimal makespan T^* is at least $\frac{1}{m} \sum_j t_j$
- Proof:
 - The total processing time is $\sum_j t_j$ min.
 - One of the m machines must have an at-least-average makespan $\frac{1}{m} \sum_j t_j$

PROOF OF 2-APPROXIMATION

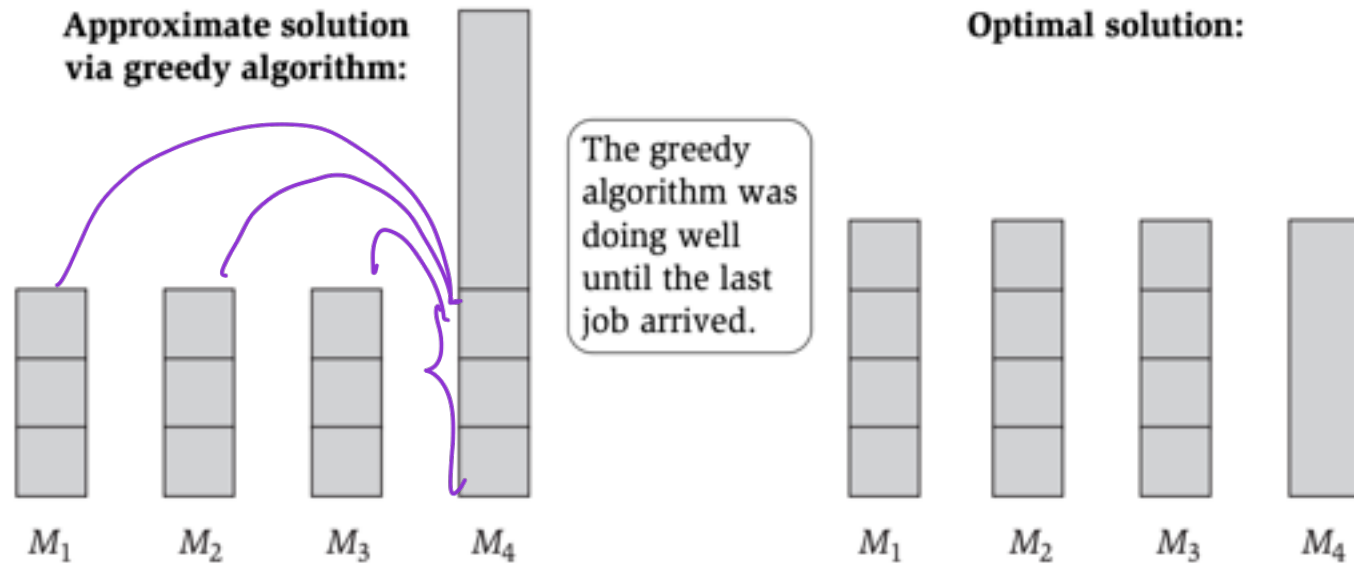
- Let L_i be the load of machine M_i at the end of the algorithm
- Let z be the last job we added to M_i
- When we added z , M_i had the smallest load at that point
 - $L_i - t_z$ — total load of the machine at t_n .
- $L_i - t_z \leq L_k$, for all other machines $k \neq i$ — all machines will be loaded by at least L_i
- Sum up over all k and divide by m :
 - $\frac{1}{m} \sum_k L_k \geq L_i - t_z$
 - t_z — time to do job z .
 - since load can only increase.
- But, $\sum_k L_k = \sum_j t_j$, the total load in the system
- By Lemma 2: $L_i - t_z \leq T^* \leq 2T^*$
 - T^* — total time.
- Overall, $L_i = t_z + (L_i - t_z) \leq 2T^*$
 - $t_z \leq \max_i t_i \leq T^*$

TIGHT ANALYSIS

- Is the factor of 2 tight?
 - Does there exist an instance in which we output a solution exactly 2 times worse than the optimal one?
- Tight example:
 - m machines, $n = m(m - 1) + 1$ jobs
 - The first $m(m - 1)$ take time $t_j = 1$
 - The last job takes time $t_n = m$
 - OPT saves a machine for the last job
 - Greedy evenly balances the first $n - 1$ jobs, and then has to add the large job to an already loaded machine

take largest last one
in more instances
forces to use 2m

TIGHT ANALYSIS



LOAD BALANCING: CAN WE DO BETTER?

- What went wrong in our tight example?
- We processed the largest job last!
 - Perhaps if we start from the large jobs we can get a better bound for the “damage” the smaller jobs can do
- New algorithm: Sorted-Balance
 1. Sort jobs in decreasing time $t_1 \geq \dots \geq t_n$
 2. Run Greedy on the sorted jobs:
 - For $j = 1 \dots n$:
 - $M_i = \operatorname{argmin}_k T_k$
 - Assign job j to machine M_i
 - $T_i += t_j$



SORTED-BALANCE ANALYSIS

- If we have fewer than m jobs, then obviously Greedy and Sorted-Balance are optimal
- Claim: If $n > m$, then $T^* \geq 2t_{m+1}$
- Proof:
 - The first m jobs take time at least t_{m+1}
 - There exists a machine assigned at least two of the first $m + 1$ jobs
 - That machine will have makespan at least $2t_{m+1}$

$$\frac{T^*}{2} \geq t_{m+1} \quad \textcircled{1}$$

SORTED-BALANCE ANALYSIS

- Theorem: Sorted-Balance is a $3/2$ approximation algorithm
- Proof:
 - Let M_i be the machine with the maximum load
 - Assume this machine has at least two jobs, and let j be the last job assigned to it
 - $j \geq m + 1$, since M_i has two jobs
 - $L_i = t_j + (L_i - t_j) \leq t_{m+1} + T^* \leq \frac{3}{2}T^*$
 - If M_i has one job, then our schedule is optimal (since the optimal schedule would also need to put the single job in M_i somewhere)

SORTED-BALANCE

- Is the factor $3/2$ tight?
- No!
- Theorem (Graham 1969): Sorted-Balance is a $4/3$ approximation
- Is the factor $4/3$ tight?
- Yes

SORTED-BALANCE

- $4/3$ example
- m machines, $n = 2m + 1$ jobs
- 2 jobs of length m , 2 jobs of length $m + 1, \dots, 2$ jobs of length $2m - 1$
- 1 job of length m
- One machine gets 3 jobs (by pigeonhole)
- OPT: $3m$
- Sorted-Balance:
 - All but one of the machines get two jobs, with processing time $3m - 1$
 - One machine gets an extra job of processing time m

WEIGHTED SET COVER

- Input:
 - A set U of n elements
 - A list S_1, \dots, S_m of subsets of U , with associated weights w_1, \dots, w_m
- Output:
 - A collection of these subsets that covers U and has **minimum** weight

WEIGHTED SET COVER

- Greedy approach: built up the solution one set at a time
- Which are good sets to include?
- Smaller weight is better
- More elements covered is better
- Pick by $w_i/|S_i|$
- Almost...
- Pick by $\frac{w_i}{|S_i \cap R|}$, where R is the set of remaining elements (as the algorithm progresses)

WEIGHTED SET COVER

- Greedy:
 - Start with $R = U$
 - While $R \neq \emptyset$:
 - Pick S_i that minimizes $w_i/|R \cap S_i|$
 - Let $c_s = w_i/|R \cap S_i|$ for all $s \in R \cap S_i$ (we'll need this line for the analysis)
 - Delete S_i from R

WEIGHTED SET COVER - ANALYSIS

- Key question in approximation algorithms: what is our benchmark for OPT?
- Typically, we do not know clean expressions for OPT (otherwise, why would we need suboptimal algorithms in the first place?)
- For weighted set cover, what should we lower bound (it's a minimization problem) OPT by?
- On the flip side, what should we upper bound the performance of our algorithm by?

WEIGHTED SET COVER - ANALYSIS

- Observation: If \mathcal{C} is the set obtained by greedy, then $\sum_{S_i \in \mathcal{C}} w_i = \sum_{s \in U} c_s$

WEIGHTED SET COVER - ANALYSIS

- Lemma 1: For every set S_k , $\sum_{S \in S_k} c_S$ is at most $w_k \cdot H_{|S_k|}$
 - $H_n = \sum_{i=1}^n 1/i$ is the n -th harmonic number
- Proof:
 - $S_k = \{s_1, s_2, \dots, s_d\}$ without loss of generality (and labeled in ordered in which they are assigned a cost)
 - s_j is covered by Greedy at some iteration
 - We know that $s_{j+1}, \dots, s_d \in R$ in that iteration
 - Thus $|S_k \cap R| \geq d - j + 1$
 - $\frac{w_k}{|S_k \cap R|} \leq \frac{w_k}{d - j + 1}$
 - Greedy selects the set with minimum average cost S_i
 - $c_{S_j} = \frac{w_i}{|S_i \cap R|} \leq \frac{w_k}{|S_k \cap R|} \leq \frac{w_k}{d - j + 1}$
 - $\sum_{S \in S_k} c_S \leq w_k \cdot \sum_{j=1}^d \frac{1}{d - j + 1} = w_k \cdot \sum_{j=1}^d \frac{1}{j} = w_k \cdot H_d$

WEIGHTED SET COVER - ANALYSIS

- Theorem: Greedy is a H_{d^*} approximation to OPT, where $d^* = \max_i |S_i|$

- Proof:

- Let \mathcal{C}^* be the optimal set cover
- $w^* = \sum_{S_i \in \mathcal{C}^*} w_i$
- For each set S_i Lemma 1 implies


$$w_i \geq \frac{\sum_{s \in S_i} c_s}{H_{|S_i|}} \geq \frac{\sum_{s \in S_i} c_s}{H_{d^*}}$$

- $\sum_{S_i \in \mathcal{C}^*} \sum_{s \in S_i} c_s \geq \sum_{s \in U} c_s$, since \mathcal{C}^* is a cover
- Putting everything together

$$w^* \geq \sum_{S_i \in \mathcal{C}^*} \frac{\sum_{s \in S_i} c_s}{H_{d^*}} \geq \frac{\sum_{s \in U} c_s}{H_{d^*}} = \frac{\sum_{S_i \in \mathcal{C}^*} w_i}{H_{d^*}}$$

WEIGHTED SET COVER

- Is this tight?
- Yes!
- Example:
 - Two tall columns of $\frac{n}{2}$ elements each
 - Two sets with weight $1 + \epsilon$ that can cover them
 - $O(\log n)$ sets of size $\frac{n}{2}, \frac{n}{4}, \dots$ with weight 1



Two sets can be used to cover everything, but the greedy algorithm doesn't find them.

WEIGHTED SET COVER

- Can we do better?
- Probably not...
- Unless $P = NP$, there is no $o(\log(n))$, i.e. better than order $\log(n)$, approximation algorithm that runs in polytime

VERTEX COVER

- Input: An undirected graph $G = (V, E)$
- Output: The smallest vertex cover, i.e. the smallest subset of vertices S such that for all $e = (u, v) \in E$, either $u \in S$ or $v \in S$
- VERTEX COVER is a special case of SET COVER
 - $U = E, S_i = \{u, v\}$ if $(u, v) \in E$
- Immediately gives a $O(\log(n))$ approximation for VERTEX COVER
 - Repeatedly delete the vertex with the largest degree and include it in the vertex cover

VERTEX COVER

- Can we do better?
- Yes!
- Repeatedly include both endpoints of an edge!
- Greedy:
 - $C = \emptyset$
 - While there exists an edge $e = (u, v)$ such that $u \notin C$ and $v \notin C$:
 - $C = C \cup \{u, v\}$

VERTEX COVER-ANALYSIS

- Observation: **Any** vertex cover is at least as large as **any** matching in G
 - A vertex cover has at least one vertex for every edge
 - In a matching, a vertex touches at most one edge
- Therefore, $OPT \geq |M|$, where OPT is the minimum vertex cover, and M is any matching
- Greedy selects a matching M_C !
 - It just puts both endpoints in
- $|C| = 2|M_C| \leq 2OPT$

VERTEX COVER-ANALYSIS

- Can we do better?
- Probably not!
- If UGC (the Unique Games Conjecture) is true, then then minimum vertex cover cannot be approximated within any constant factor better than 2.

SUMMARY

- LOAD BALANCING (11.1)
- SET COVER (11.3)
- VERTEX COVER (DPV 9.2.1)